

Making
Cheetah
Faster



What is Cheetah?

- In-Memory Database (Java)
- Developed by *Alan Lu*
- JSON ~> Table
- Exploration of various *data stores*
- Queried via SQL statements
- *Single-threaded...*

Why Cheetah?

- Faster than *disk-optimized* databases (~6x)
- Eliminates *seek time* when querying...
- *Faster* and *more predictable* performance
- Critical for *real-time* response

Overview: Cheetah

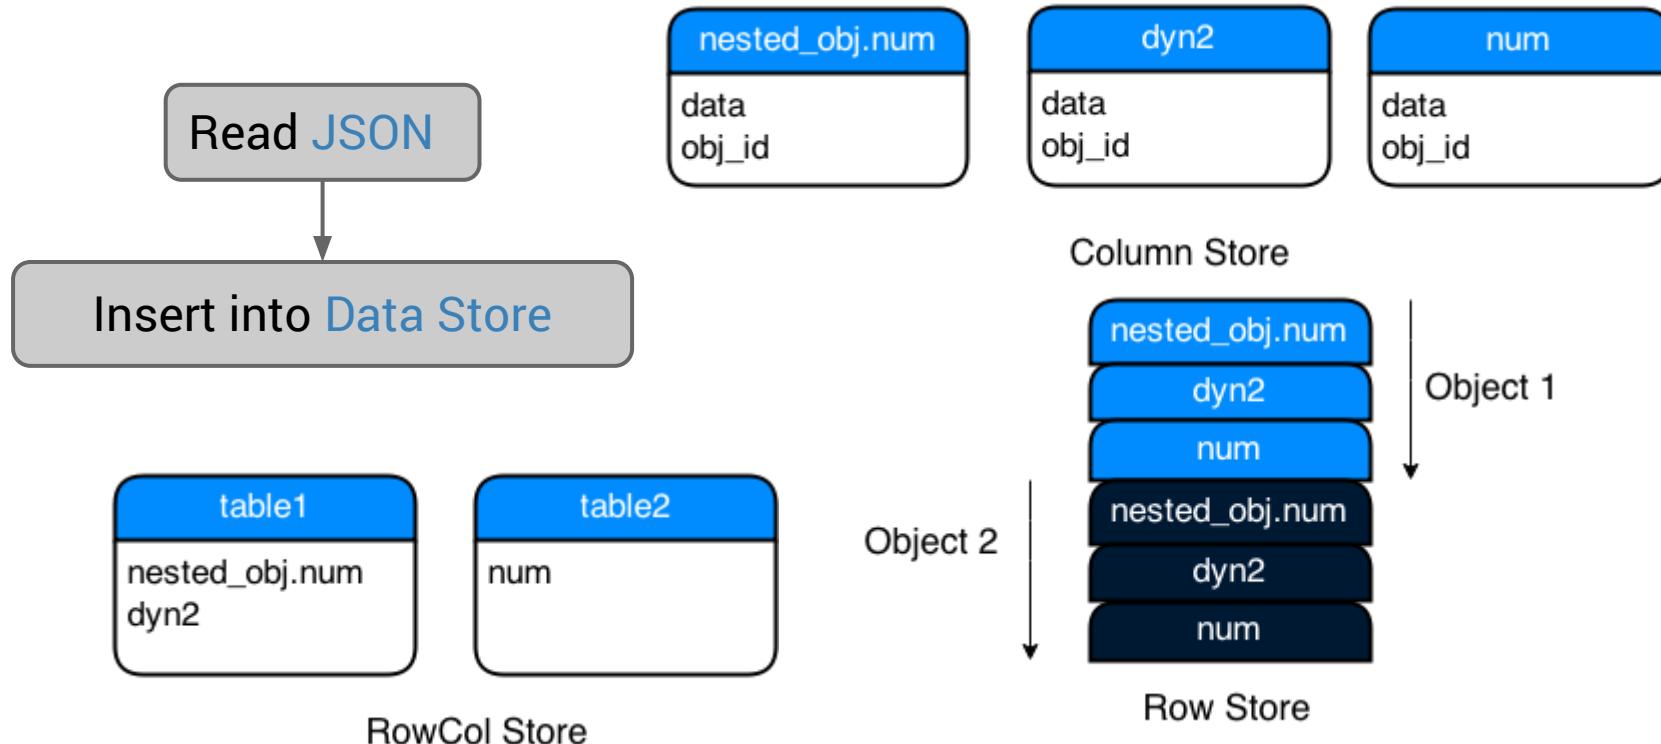
- Current *Architecture*
- Current *Data Stores*
- Exploring the Column Store Implementation
- Speeding Up *Queries* and *Insertion!*

Cheetah: Current Architecture

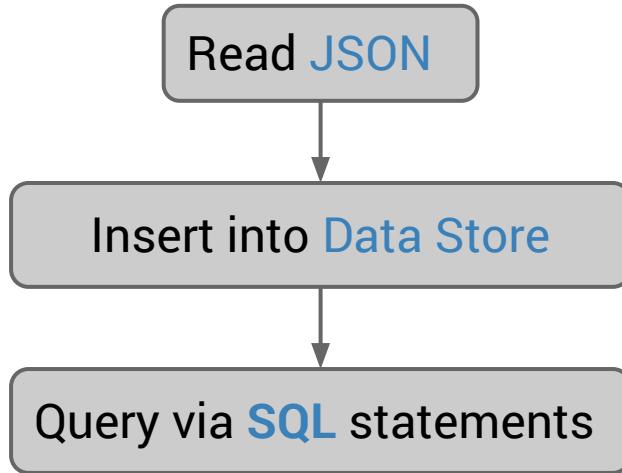
Read JSON

```
[  
  {  
    nested_obj: {  
      num: 4507  
    },  
    dyn2: true,  
    nested_arr: [  
      "checking",  
      "here"  
    ],  
    str2: "GB",  
    num: 9507  
  },  
  {}  
]
```

Cheetah: Current Architecture



Cheetah: Current Architecture



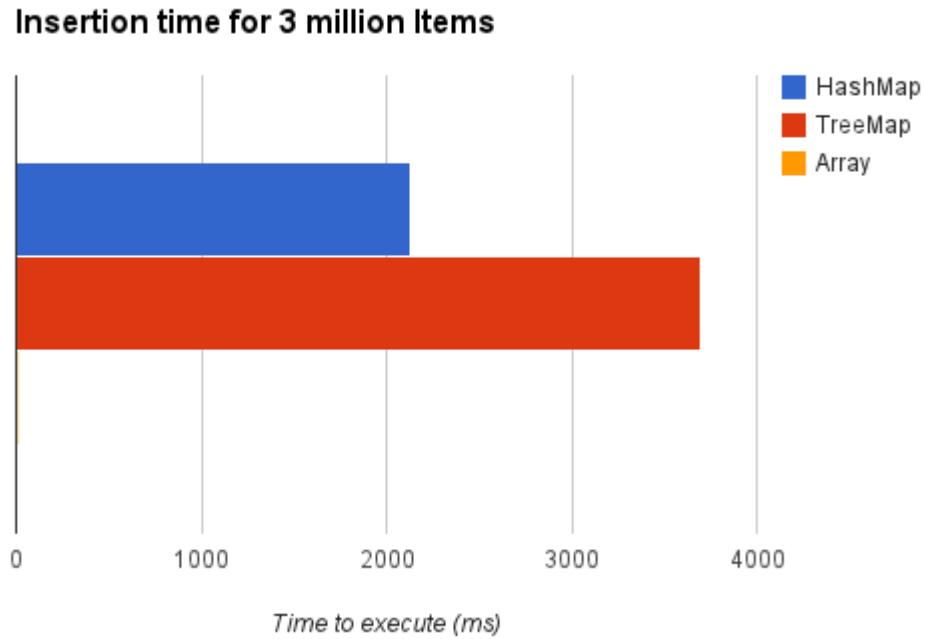
```
SELECT dyn2  
WHERE num  
BETWEEN 0 AND 10
```

Cheetah: Current Data stores

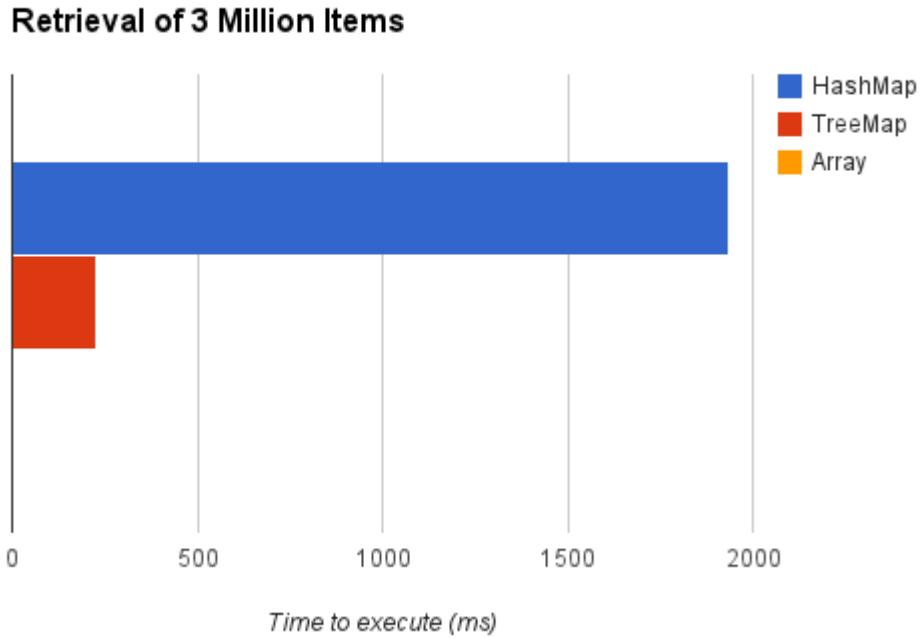


- Δ Each *key* has a table
- Δ No *Schema* required
- Δ Uses two arrays per table
- Δ *Excellent* for sparse data
- Δ *Not so great* for *Querying*

Cheetah: Other Data Structures?



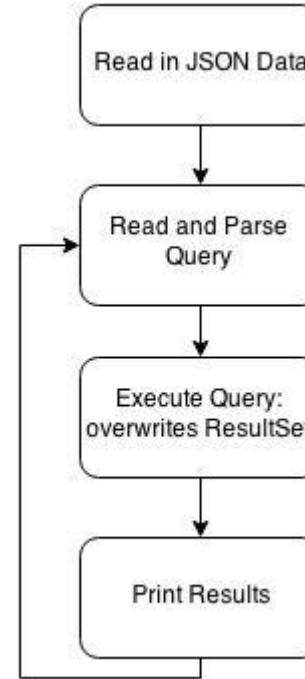
Cheetah: Other Data Structures?



Cheetah: Running Queries in Parallel

Original program flow:

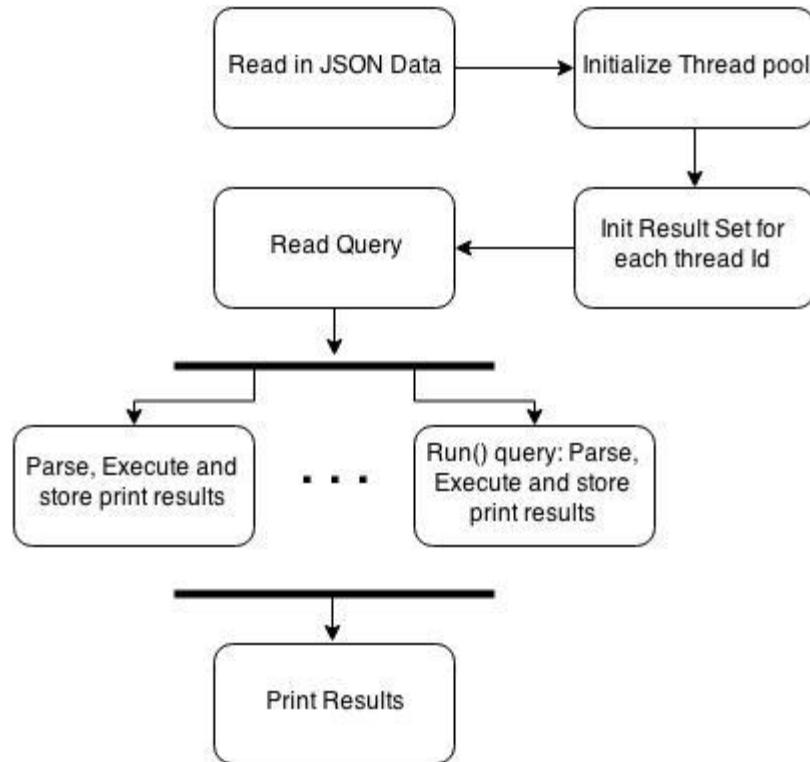
- 1) Read in database
- 2) Read and perform queries one by one
 - a) Parse query
 - b) Execute query → writes to a result set
 - c) Print result information
- 3) Next query clears and overwrites result set



Cheetah: Running Queries in Parallel

Parallel program flow:

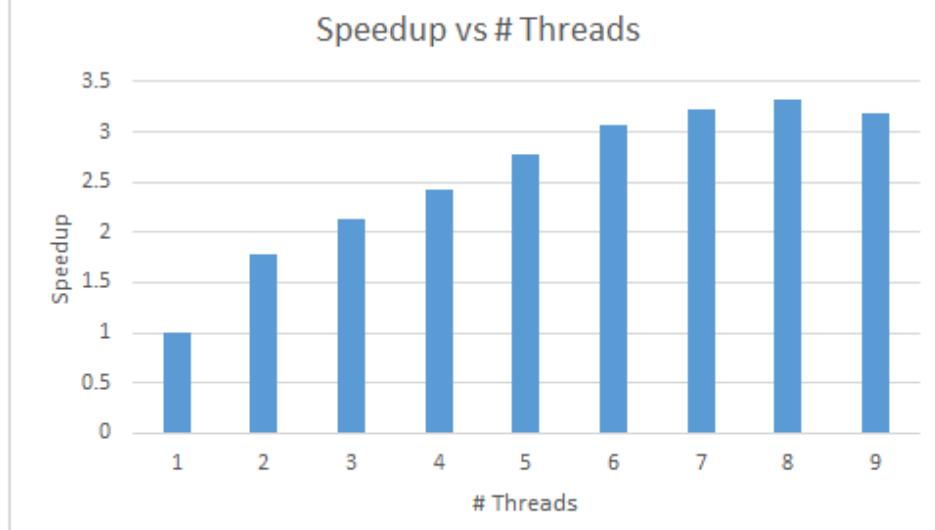
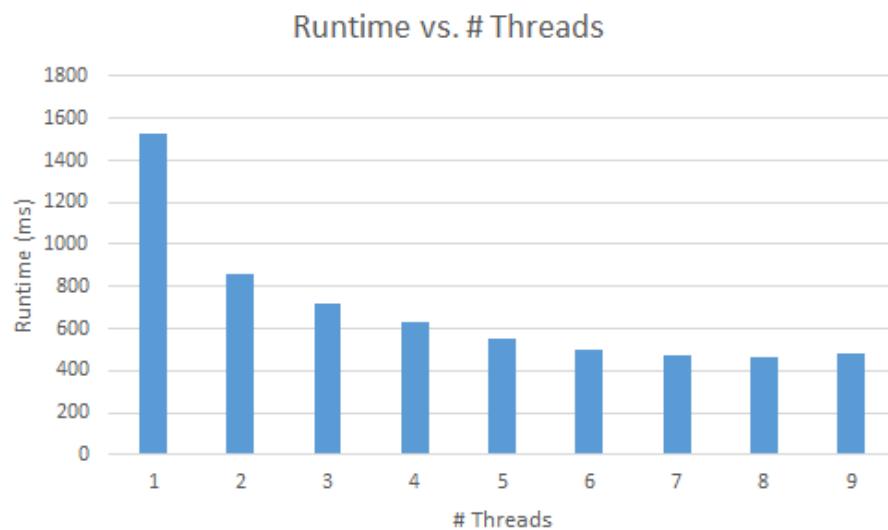
- 1) Read in database
- 2) Initialize thread pool
- 3) Initialize result set: each thread has it's own result set
- 4) Run queries in parallel
- 5) Print results



Cheetah: Issues with parallelization

- Results Sets are large: 3 large arrays in each result set
- Lack of memory

Cheetah: Running Queries in Parallel



Cheetah: Java Profiling with YourKit

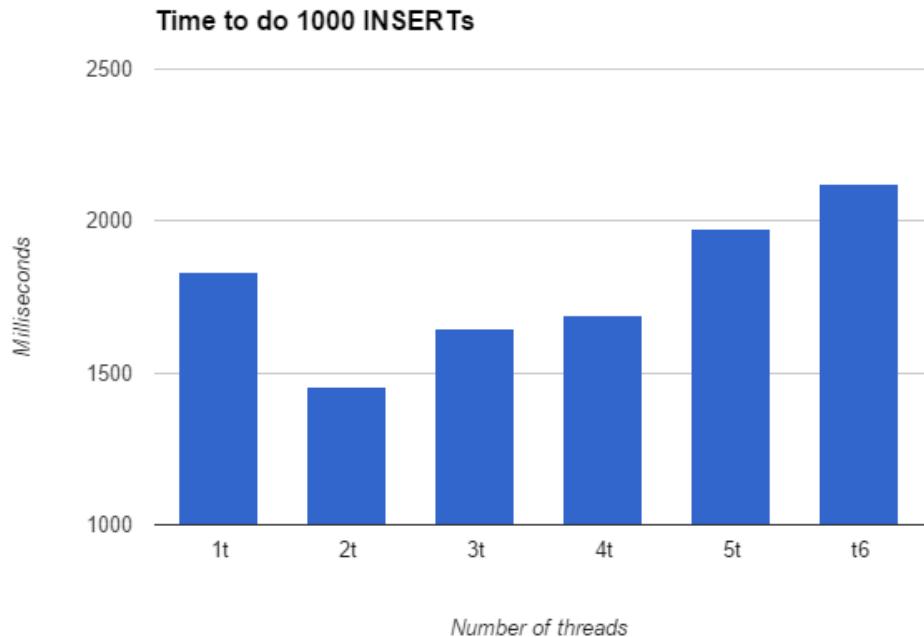
Sequential Execution time: ~7000ms

Threadpool with 1 thread time: ~1500ms

	% of time for Original program (Sequential)	% of time for parallelized program (with 1 thread)
Executing Querying	9%	~0%
Printing Summary	47%	15%

Cheetah: Faster Insertions

- Parallel inserts of JSON objects
- Coarse grained lock
- Fine grained lock works better for sparse data



Cheetah: What's next?

- Exclusion SQL statement (e.g NOT IN)
- Explore other data store configurations
- Cache hit/miss mechanism
- Explore lock granularities: table, column, hybrid

Cheetah: Conclusion

- Explored different data structures
- Added functionality for **SQL INSERT** statement
- Achieved speedup by adding multi-threaded support for read and write



Thanks for listening!